

Supporting Information for BioInteract: Interpretable Drug–Target Interaction Prediction via Residue-Level Cross-Attention with Biological Prior Knowledge

This document reports only values recoverable from the released Davis records, configuration files, saved result files, and source code included with this revision. “Classifier score” denotes the sigmoid output for the binary high-affinity class; it is not an experimental affinity. Model-native attention and Grad-CAM values are attribution scores, not contact assignments.

Model configuration and input features

Table S1: Model architecture confirmed from the checkpoint-embedded model configurations.

Component	Archived setting
Drug encoder	GINE; 3 layers; hidden dimension 256; dropout 0.20
Atom and bond input	52 node features; 16 edge features
Protein language model	<code>esm2_t30_150M_UR50D</code> ; 640-dimensional residues
Protein projection	256-dimensional projection; 32-dimensional domain-label embedding
Cross-attention	8 heads; attention dimension 256; dropout 0.10
Pooling	Gated
Prediction head	Hidden dimensions 256 and 128; dropout 0.30; classification task

Table S2: Protein-residue input channels and their availability.

Channel	Archived implementation and limitation
ESM-2	Frozen 640-dimensional residue embeddings; proteins are truncated at 1,200 residues.
Physicochemical vector	Four values per residue: hydrophobicity, normalized molecular weight, charge, and polarity.
Domain-label channel	The configuration reserves a 32-dimensional embedding. The label builder uses <code>NONE</code> when no per-protein JSON annotation is available.
Annotation availability	No <code>data/domain_annotations/*.json</code> files are present in the released package. Thus the reproducible Davis input uses the unknown label at every residue and is not curated Pfam/InterPro information.

Table S3: Training provenance for the historical entity-split checkpoints.

Setting	Value
Checkpoint-embedded evidence	Model architecture only; the three entity-split checkpoints do not retain optimizer, scheduler, warmup, or loss state
Random and Drug-ID logs	Final-run headers record weighted BCE, label smoothing 0.05, AdamW (10^{-4} , 10^{-5}), and cosine warmup
Target-ID log matched to its checkpoint	Older <code>run_split.py</code> header; only loss and validation AUROC are logged, so optimizer, warmup, and smoothing are not recoverable from the artifact
Shared data provenance	Davis binary high-affinity task; seed 42; archived 0.70/0.10/0.20 split proportions
Current generic workflow	<code>src.cli.train</code> is a new-run entrypoint, not historical provenance for the released checkpoints
Threshold protocol	Validation F1 threshold is frozen before each canonical test evaluation

Table S4: Computational environment documented during final reproducibility audit.

Item	Value
Python	3.10.14
PyTorch	2.4.0+cu121
PyTorch Geometric	2.6.1
RDKit	2025.3.6
<code>fair-esm</code>	2.0.0
Biopython and RapidFuzz	1.87 and 3.14.5
CUDA availability during audit	Available
Audit GPU	NVIDIA GeForce RTX 4080
Strict-split run record	<code>device: cuda</code> in <code>strict_split_metrics.json</code>
Checkpoint files available	<code>best.pt</code> , <code>best_random.pt</code> , <code>best_cold_target.pt</code> , <code>best_cold_drug.pt</code>

Dataset, partitions, and metrics

Table S5: Canonical current-release Davis test-partition summaries. Values are from the full-precision checkpoint reconstruction, frozen validation thresholds, and the associated prediction files.

Partition	Test pairs	Positives	AUROC	AUPRC	F1	Precision	Recall
Random	6,011	276	0.904	0.560	0.565	0.602	0.533
Target-ID-held-out	5,984	250	0.930	0.525	0.534	0.517	0.552
Drug-ID-held-out	5,746	245	0.733	0.167	0.100	0.273	0.061

Table S6: Bootstrap intervals for the canonical current-release entity-based partitions. Each interval uses 600 ordinary pair-level bootstrap replicates sampled with replacement at the test-pair level, NumPy seed 42, and the 2.5th and 97.5th percentiles of valid replicate metrics. Replicates containing only one class are discarded. The main figure shows point estimates only.

Partition	AUROC 95% interval	AUPRC 95% interval
Random	[0.881, 0.924]	[0.498, 0.624]
Target-ID-held-out	[0.911, 0.948]	[0.454, 0.584]
Drug-ID-held-out	[0.703, 0.763]	[0.129, 0.211]

Table S7: Exact-sequence-grouped strict within-Davis evaluations from `strict_split_metrics.json`. Threshold-dependent quantities use the corresponding validation-set threshold.

Pair accounting					
Protocol	Train pairs	Validation pairs	Test pairs	Test positives	Excluded cross-block pairs
Exact-sequence-grouped cold-target	21,080	2,992	5,984	208	0
Exact-sequence-grouped cold-both	15,190	264	1,144	38	13,458

Test metrics							
Protocol	AUROC	AUPRC	Threshold	F1	Precision	Recall	
Exact-sequence-grouped cold-target	0.873	0.383	0.838	0.468	0.517	0.428	
Exact-sequence-grouped cold-both	0.552	0.038	0.936	0.000	0.000	0.000	

For the cold-both protocol, drug identifiers and exact target-sequence groups are partitioned independently at 70/10/20. Only pairs in the corresponding train–train, validation–validation, and test–test blocks are retained; the remaining cross-block pairs are excluded. The cold-both validation block contains 10 positive pairs.

Table S8: Source and protocol context for baseline values reproduced from the comparison table in the originally submitted manuscript. The table identifies the numerical source, reported evaluation setting, protein-representation context, and the appropriate descriptive interpretation of the comparison.

Method	Numerical source	Original comparison protocol	Binary criterion	Protein representation	Interpretation in this revision
DeepDTA	Original submitted comparison table	Random, Target-ID-held-out, and Drug-ID-held-out	30 nM Davis affinity criterion	Baseline-native representation; no matched frozen ESM-2 input	Descriptive contextual comparison; not a matched-pretraining or common-code benchmark
GraphDTA	Original submitted comparison table	Random, Target-ID-held-out, and Drug-ID-held-out	30 nM Davis affinity criterion	Baseline-native representation; no matched frozen ESM-2 input	Descriptive contextual comparison; not a matched-pretraining or common-code benchmark
AttentionDTA	Original submitted comparison table	Random, Target-ID-held-out, and Drug-ID-held-out	30 nM Davis affinity criterion	Baseline-native representation; no matched frozen ESM-2 input	Descriptive contextual comparison; not a matched-pretraining or common-code benchmark
✈ MolTrans	Original submitted comparison table	Random, Target-ID-held-out, and Drug-ID-held-out	30 nM Davis affinity criterion	Baseline-native representation; no matched frozen ESM-2 input	Descriptive contextual comparison; not a matched-pretraining or common-code benchmark
TransformerCPI	Original submitted comparison table	Random, Target-ID-held-out, and Drug-ID-held-out	30 nM Davis affinity criterion	Baseline-native representation; no matched frozen ESM-2 input	Descriptive contextual comparison; not a matched-pretraining or common-code benchmark
DrugBAN	Original submitted comparison table	Random, Target-ID-held-out, and Drug-ID-held-out	30 nM Davis affinity criterion	Baseline-native representation; no matched frozen ESM-2 input	Descriptive contextual comparison; not a matched-pretraining or common-code benchmark
BioInteract	Archived prediction CSVs and result JSONs	Random, Target-ID-held-out, and Drug-ID-held-out	30 nM Davis affinity criterion	Frozen ESM-2 150M residue embeddings	Traceable within BioInteract; not a cross-method matched-encoder comparison

Reproducibility details for revision analyses

Identifier-partition similarity audits. Drug similarities use RDKit 2025.3.6 Morgan fingerprints (radius 2, 1,024 bits, `includeChirality=False`) and Tanimoto similarity. Protein similarities for the original Target-ID-held-out audit use RapidFuzz 3.14.5 `fuzz.ratio/100`. For sequences x, y , this is the normalised Indel similarity $1 - d_{\text{Indel}}/(|x| + |y|)$, where substitutions are represented as one deletion plus one insertion.

Strict-test global alignment. Biopython 1.87 `Align.PairwiseAligner` is used in global mode with `match = 1`, `mismatch = 0`, `gap-opening = -1`, and `gap-extension = -0.5`. Full-length identity is the number of exact-residue matches divided by alignment length; the maximum across all training targets is retained for each held-out target.

Bootstrap intervals. The saved 95% intervals use ordinary, unstratified pair-level resampling with replacement, 600 requested replicates, and NumPy random seed 42. Each replicate contains the same number of test pairs as the corresponding test set; one-class replicates are skipped. Bounds are the 2.5th and 97.5th percentiles of the valid AUROC or AUPRC replicate distribution.

Grad-CAM. Hooks are attached to `drug_encoder.layers.2`, the third and final GINE message-passing layer. Gradients are taken from the pre-sigmoid positive-class logit, pooled across graph nodes to obtain channel weights, rectified, and then normalised separately within each molecular graph.

Attention sparsity. For pair q , residue scores $s_{qj} = \sum_i M_{qij}$ are normalised as $\tilde{s}_{qj} = s_{qj} / \max_k s_{qk}$. The reported fraction is $|\{(q, j) : \tilde{s}_{qj} < 0.1\}| / \sum_q L_q$. It uses all 1,506 positive Davis pairs with the archived `checkpoints/best.pt`, not a single test set. Every pair is normalised before the 1,301,380 valid residue scores are concatenated; thus original training, validation, and test memberships are all included and longer targets contribute more residue-level observations. The archived flat-score array is `BioInteract/results/figure_data/attention_distribution.npz`; it is the sole source of the Figure 5 distribution and cumulative curve.

Exact summary of the archived pair-normalised residue-attribution array used for

Figure 5.	Quantity	Value
	Positive Davis pairs	1,506
	Residue-score observations	1,301,380
	Mean / standard deviation	0.004351 / 0.046243
	Median	0.000272
	95th / 99th percentile	0.003617 / 0.057193
	Fraction below 0.1	0.992319
	Attribution mass in top 1% / 5% scores	0.788528 / 0.904465

Exact-sequence-grouped target diagnostic controls. All values use the 5,984-pair strict target test set with 208 positives.

Control	AUROC	AUPRC
BioInteract	0.873	0.383
Drug-only Morgan logistic regression	0.764	0.164
Drug-promiscuity prior	0.764	0.163
Target-only constant prior	0.500	0.035
Label-shuffled drug-only control	0.407	0.029

Feature and input audits

Table S9: Atom and bond feature specification from `src/data/mol_graph.py`.

Level	Feature family	Dimensions
Atom	Atom type one-hot (20 choices including unknown)	20
Atom	Hybridization, formal charge, and hydrogen-count one-hot vectors	18
Atom	Degree, aromaticity, and ring membership	3
Atom	Ring-size membership for sizes 3–8	6
Atom	Donor/acceptor, Crippen logP and MR, Gasteiger charge	5
Atom total		52
Bond	Bond type, conjugation, ring membership, stereo, and direction	16

Table S10: Davis nominal ABL1 input audit. Every listed record has the same complete SHA-256 digest `cd6a78b82fede215cdb80a797f228098a41178fdde0753728f3781e6221fb1ee`; values in the last columns are the archived D0010 and D0017 affinities in nM. The nominal labels do not alter the model sequence input.

Target ID	Nominal name	Length	Matches ABL1	D0010	D0017
T0001	ABL1(E255K)	1167	Yes	0.047	0.047
T0002	ABL1(F317I)	1167	Yes	0.63	0.10
T0003	ABL1(F317I)p	1167	Yes	0.18	0.041
T0004	ABL1(F317L)	1167	Yes	0.11	0.032
T0005	ABL1(F317L)p	1167	Yes	0.029	0.019
T0006	ABL1(H396P)	1167	Yes	0.062	0.025
T0007	ABL1(H396P)p	1167	Yes	0.057	0.046
T0008	ABL1(M351T)	1167	Yes	0.037	0.016
T0009	ABL1(Q252H)	1167	Yes	0.086	0.037
T0010	ABL1(Q252H)p	1167	Yes	0.039	0.064
T0011	ABL1(T315I)	1167	Yes	21	890
T0012	ABL1(T315I)p	1167	Yes	3.6	120
T0013	ABL1(Y253F)	1167	Yes	0.036	0.058
T0014	ABL1	1167	Yes	0.12	0.029
T0015	ABL1p	1167	Yes	0.057	0.046

Table S11: Saved non-mutant EGFR–D0014 functional-group attribution values.

Functional group	Attribution
Carbonyl	0.431
Amide	0.389
Ether	0.136
Halogen	0.134
Amino	0.090
Aromatic ring	0.000
Heterocycle N	0.000

Table S12: Saved non-mutant BRAF–D0023 failure-mode summary. The archived affinity is positive at the study threshold, whereas the classifier score is low. The score and attributions are model outputs, not an explanation of binding.

Quantity	Value
High-affinity-class score	0.099
Archived Davis affinity (nM)	0.19
Aromatic-ring attribution	0.059
Hydroxyl attribution	0.000
Heterocycle-N attribution	0.000

Table S13: Trainable parameter count recovered from `interpretability_report.json`.

Quantity	Value
Total trainable parameters	2,442,083